

Import Library Files

```
from google.colab import drive
```

```
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from sklearn.compose import ColumnTransformer
```

```
from sklearn.preprocessing import OneHotEncoder
```

```
from sklearn.preprocessing import StandardScaler
```

2. Load the dataset

```
data = pd.read_csv('/content/drive/MyDrive/Churn_Modelling.csv')
```

3.1 UNIVARIATE ANALYSIS

```
l=['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary']
```

```
for i in l:
```

```
    sns.displot(data=data[i],kde=True)
```

```
l=['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'EstimatedSalary']
```

```
fig, (ax1, ax2, ax3, ax4, ax5, ax6) = plt.subplots(nrows=6, ncols=1, figsize=(10,20))
```

```
data.boxplot(column=l[0],grid=False,color='blue',ax=ax1)
```

```
data.boxplot(column=l[1],grid=False,color='blue',ax = ax2)
```

```

data.boxplot(column=I[2],grid=False,color='blue',ax = ax3)
data.boxplot(column=I[3],grid=False,color='blue',ax = ax4)
data.boxplot(column=I[4],grid=False,color='blue',ax = ax5)
data.boxplot(column=I[5],grid=False,color='blue',ax = ax6)
plt.tight_layout()

```

```

import warnings
warnings.filterwarnings("ignore")
fig, (ax1, ax2, ax3) = plt.subplots(nrows=3, ncols=1, figsize=(16,16))
sns.countplot(data.HasCrCard,ax=ax1)
sns.countplot(data.IsActiveMember,ax=ax2)
sns.countplot(data.Exited,ax=ax3)
plt.tight_layout()

```

3.2 BI - VARIATE ANALYSIS

```

for i in range(len(I)-1):
    for j in range(i+1,len(I)):
        sns.relplot(x = I[i],y = I[j],data = data)

```

3.3 MULTI - VARIATE ANALYSIS

```
sns.catplot(x='Gender', y='Age', hue='HasCrCard', data=data)
```

<seaborn.axisgrid.FacetGrid at 0x7fdd89a27210>

```
sns.pairplot(data = data,hue='Exited')
```

<seaborn.axisgrid.PairGrid at 0x7fdd87c0c290>

4. Perform descriptive statistics on the dataset

```
data.head()
```

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure			
	Balance	NumOfProducts	HasCrCard	IsActiveMember		EstimatedSalary	Exited			
1	2	15647311	645	608	2	0	41	1	83807.86	1
	0	1	112542.58	0						
5	6	15574012	302	645	2	1	44	8	113755.78	2
	1	0	149756.71	1						
10	11	15767821	109	528	0	1	31	6	102016.72	2
	0	0	80181.12	0						
15	16	15643966	561	616	1	1	45	3	143129.41	2
	0	1	64327.26	0						
26	27	15736816	1605	756	1	1	36	2	136815.64	1
	1	1	170041.95	0						

```
data.describe()
```

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure
	Balance	NumOfProducts	HasCrCard	IsActiveMember		EstimatedSalary	Exited
count	3354.000000	3.354000e+03	3354.000000	3354.000000	3354.000000	3354.000000	3354.000000
	3354.000000	3354.000000	3354.000000	3354.000000	3354.000000	3354.000000	3354.000000
	3354.000000	3354.000000					

mean	4993.122242	1.568997e+07	834.928145	651.885808	0.798151	0.551580
	38.594812	4.960644	111127.251270	1.386106	0.705426	0.496720
	101173.425200	0.237627				
std	2889.712337	7.220517e+04	467.984617	66.341508	0.744870	0.497407
	6.171482	2.910065	23930.791436	0.579239	0.455919	0.500064
	57475.269109	0.425693				
min	2.000000	1.556571e+07	0.000000	522.000000	0.000000	0.000000
	29.000000	0.000000	3768.690000	1.000000	0.000000	0.000000
	11.580000	0.000000				
25%	2469.250000	1.562732e+07	443.250000	601.000000	0.000000	0.000000
	34.000000	2.000000	96579.825000	1.000000	0.000000	0.000000
	53183.340000	0.000000				
50%	4986.500000	1.568912e+07	847.500000	652.000000	1.000000	1.000000
	38.000000	5.000000	113904.805000	1.000000	1.000000	0.000000
	101348.755000	0.000000				
75%	7483.750000	1.575380e+07	1250.000000	705.000000	1.000000	1.000000
	43.000000	7.000000	129621.140000	2.000000	1.000000	1.000000
	150202.787500	0.000000				
max	9999.000000	1.581569e+07	1628.000000	777.000000	2.000000	1.000000
	53.000000	10.000000	149238.970000	4.000000	1.000000	1.000000
	199970.740000	1.000000				

data.dtypes

RowNumber	int64
CustomerId	int64
Surname	int64
CreditScore	int64
Geography	int64
Gender	int64
Age	int64
Tenure	int64
Balance	float64
NumOfProducts	int64
HasCrCard	int64
IsActiveMember	int64

EstimatedSalary float64

Exited int64

dtype: object

data.skew()

RowNumber 0.011381

CustomerId 0.010861

Surname -0.033943

CreditScore -0.024726

Geography 0.344510

Gender -0.207520

Age 0.415624

Tenure 0.034787

Balance -0.691529

NumOfProducts 1.450814

HasCrCard -0.901691

IsActiveMember 0.013125

EstimatedSalary 0.002697

Exited 1.233424

dtype: float64

5. Handle the Missing values.

data.isnull().any()

RowNumber False

CustomerId False

Surname False

CreditScore False

Geography False

Gender False

Age False

Tenure False

```
Balance      False
NumOfProducts  False
HasCrCard     False
IsActiveMember  False
EstimatedSalary  False
Exited        False
dtype: bool
No missing values
```

6. Find the outliers and replace the outliers

```
data['CreditScore'].describe()
```

```
count    10000.000000
mean      650.528800
std       96.653299
min       350.000000
25%       584.000000
50%       652.000000
75%       718.000000
max       850.000000
```

```
Name: CreditScore, dtype: float64
```

```
data['Age'].describe()
```

```
count    10000.000000
mean      38.921800
std       10.487806
min       18.000000
25%       32.000000
50%       37.000000
75%       44.000000
max       92.000000
```

Name: Age, dtype: float64

data['Balance'].describe()

count 10000.000000

mean 76485.889288

std 62397.405202

min 0.000000

25% 0.000000

50% 97198.540000

75% 127644.240000

max 250898.090000

Name: Balance, dtype: float64

l=['Balance','Age','CreditScore']

for i in l:

percentile_least = data[i].quantile(0.1)

percentile90 = data[i].quantile(0.9)

data = data[(data[i]<percentile90)& (data[i]>percentile_least)]

data['CreditScore'].describe()

count 3354.000000

mean 651.885808

std 66.341508

min 522.000000

25% 601.000000

50% 652.000000

75% 705.000000

max 777.000000

Name: CreditScore, dtype: float64

data['Age'].describe()

count 3354.000000

mean 38.594812

```
std      6.171482
min      29.000000
25%      34.000000
50%      38.000000
75%      43.000000
max      53.000000
```

Name: Age, dtype: float64

```
data['Balance'].describe()
```

```
count    3354.000000
mean     111127.251270
std      23930.791436
min       3768.690000
25%      96579.825000
50%     113904.805000
75%     129621.140000
max     149238.970000
```

Name: Balance, dtype: float64

The STD have reduced drastically indicating removal of outliers.

7. Check for Categorical columns and perform encoding.

```
from sklearn.preprocessing import LabelEncoder
```

```
encoder=LabelEncoder()
```

```
for i in data:
```

```
    if data[i].dtype=='object':
```

```
        data[i]=encoder.fit_transform(data[i])
```

```
data.head()
```

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure
	Balance	NumOfProducts	HasCrCard	IsActiveMember		EstimatedSalary	Exited

1	2	15647311	645	608	2	0	41	1	83807.86	1
	0	1	112542.58	0						
5	6	15574012	302	645	2	1	44	8	113755.78	2
	1	0	149756.71	1						
10	11	15767821	109	528	0	1	31	6	102016.72	2
	0	0	80181.12	0						
15	16	15643966	561	616	1	1	45	3	143129.41	2
	0	1	64327.26	0						
26	27	15736816	1605	756	1	1	36	2	136815.64	1
	1	1	170041.95	0						

8. Split the data into dependent and independent variables.

data.shape

(3354, 14)

x = data.iloc[:,13]

y = data.iloc[:,13]

y.head()

1 0

5 1

10 0

15 0

26 0

Name:Exited, dtype: int64

x.head()

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure			
	Balance	NumOfProducts	HasCrCard	IsActiveMember		EstimatedSalary				
1	2	15647311	645	608	2	0	41	1	83807.86	1
	0	1	112542.58							
5	6	15574012	302	645	2	1	44	8	113755.78	2
	1	0	149756.71							
10	11	15767821	109	528	0	1	31	6	102016.72	2
	0	0	80181.12							

15	16	15643966	561	616	1	1	45	3	143129.41	2
	0	1	64327.26							
26	27	15736816	1605	756	1	1	36	2	136815.64	1
	1	1	170041.95							

9. Scale the independent variables

```
from sklearn.preprocessing import StandardScaler
```

```
sc = StandardScaler()
```

```
x = sc.fit_transform(x)
```

10. Split the data into training and testing

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2, random_state=0)
```

```
x_train.shape
```

```
(2683, 13)
```

```
y_train.shape
```

```
(2683,)
```

```
x_test.shape
```

```
(671, 13)
```

```
y_test.shape
```

```
(671,)
```